

# CS187 Project Segment 4: Semantic Interpretation – Question Answering

December 11, 2021

```
[ ]: # Please do not change this cell because some hidden tests might depend on it.  
import os
```

```
# Otter grader does not handle ! commands well, so we define and use our  
# own function to execute shell commands.
```

```
def shell(commands, warn=True):  
    """Executes the string `commands` as a sequence of shell commands.  
  
    Prints the result to stdout and returns the exit status.  
    Provides a printed warning on non-zero exit status unless `warn`  
    flag is unset.  
    """  
    file = os.popen(commands)  
    print (file.read().rstrip('\n'))  
    exit_status = file.close()  
    if warn and exit_status != None:  
        print(f"Completed with errors. Exit status: {exit_status}\n")  
    return exit_status
```

```
shell("""  
ls requirements.txt >/dev/null 2>&1  
if [ ! $? = 0 ]; then  
    rm -rf .tmp  
    git clone https://github.com/cs187-2021/project4.git .tmp  
    mv .tmp/requirements.txt ./  
    rm -rf .tmp  
fi  
pip install -q -r requirements.txt  
""")
```

```
[ ]: # Initialize Otter  
import otter  
grader = otter.Notebook()
```

The goal of semantic parsing is to convert natural language utterances to a meaning representation such as a *logical form* expression or a *SQL query*. In the previous project segment, you built a parsing system to reconstruct parse trees from the natural-language queries in the ATIS dataset. However, that only solves an intermediary task, not the end-user task of obtaining answers to the queries.

In this final project segment, you will go further, building a semantic parsing system to convert English queries to SQL queries, so that by consulting a database you will be able to answer those questions. You will implement both a rule-based approach and an end-to-end sequence-to-sequence (seq2seq) approach. Both algorithms come with their pros and cons, and by the end of this segment you should have a basic understanding of the characteristics of the two approaches.

## 0.1 Goals

1. Build a semantic parsing algorithm to convert text to SQL queries based on the syntactic parse trees from the last project.
2. Build an attention-based end-to-end seq2seq system to convert text to SQL.
3. Improve the attention-based end-to-end seq2seq system with self-attention to convert text to SQL.
4. Discuss the pros and cons of the rule-based system and the end-to-end system.
5. (Optional) Use the state-of-the-art pretrained transformers for text-to-SQL conversion.

This will be an extremely challenging project, so we recommend that you start early.

## 1 Setup

```
[ ]: import copy
import datetime
import math
import re
import sys
import warnings

import wget
import nltk
import sqlite3
import torch
import torch.nn as nn
import torchtext.legacy as tt

from cryptography.fernet import Fernet
from func_timeout import func_set_timeout
from torch.nn.utils.rnn import pack_padded_sequence as pack
from torch.nn.utils.rnn import pad_packed_sequence as unpack
from tqdm import tqdm
from transformers import BartTokenizer, BartForConditionalGeneration
```

```
[ ]: # Set random seeds
seed = 1234
torch.manual_seed(seed)
# Set timeout for executing SQL
TIMEOUT = 3 # seconds

# GPU check: Set runtime type to use GPU where available
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print (device)

[ ]: ## Download needed scripts and data
os.makedirs('data', exist_ok=True)
os.makedirs('scripts', exist_ok=True)
source_url = "https://raw.githubusercontent.com/nlp-course/data/master"

# Grammar to augment for this segment
if not os.path.isfile('data/grammar'):
    wget.download(f"{source_url}/ATIS/grammar_distrib4.crypt", out="data/")

# Decrypt the grammar file
key = b'bfksTY2BJ5VKKK9xZb1PDDLgKdu7KCDFYfVePSEfGY='
fernet = Fernet(key)
with open('./data/grammar_distrib4.crypt', 'rb') as f:
    restored = Fernet(key).decrypt(f.read())
with open('./data/grammar', 'wb') as f:
    f.write(restored)

# Download scripts and ATIS database
wget.download(f"{source_url}/scripts/trees/transform.py", out="scripts/")
wget.download(f"{source_url}/ATIS/atis_sqlite.db", out="data/")

[ ]: # Import downloaded scripts for parsing augmented grammars
sys.path.insert(1, './scripts')
import transform as xform
```

## 2 Semantically augmented grammars

In the first part of this project segment, you'll be implementing a rule-based system for semantic interpretation of sentences. Before jumping into using such a system on the ATIS dataset – we'll get to that soon enough – let's first work with some trivial examples to get things going.

The fundamental idea of rule-based semantic interpretation is the rule of compositionality, that *the meaning of a constituent is a function of the meanings of its immediate subconstituents and the syntactic rule that combined them*. This leads to an infrastructure for specifying semantic interpretation in which each syntactic rule in a grammar (in our case, a context-free grammar) is associated with a semantic rule that applies to the meanings associated with the elements on the right-hand side of the rule.

## 2.1 Example: arithmetic expressions

As a first example, let's consider an augmented grammar for arithmetic expressions, familiar from lab 3-1. We again use the function `xform.parse_augmented_grammar` to parse the augmented grammar. You can read more about it in the file `scripts/transform.py`.

```
[ ]: arithmetic_grammar, arithmetic_augmentations = xform.parse_augmented_grammar(
    """
    ## Sample grammar for arithmetic expressions

    S -> NUM                                : lambda Num: Num
      / S OP S                              : lambda S1, Op, S2: Op(S1, S2)

    OP -> ADD                                : lambda Op: Op
      / SUB
      / MULT
      / DIV

    NUM -> 'zero'                            : lambda: 0
      / 'one'                                : lambda: 1
      / 'two'                                : lambda: 2
      / 'three'                              : lambda: 3
      / 'four'                               : lambda: 4
      / 'five'                               : lambda: 5
      / 'six'                                : lambda: 6
      / 'seven'                              : lambda: 7
      / 'eight'                             : lambda: 8
      / 'nine'                              : lambda: 9
      / 'ten'                                : lambda: 10

    ADD -> 'plus' / 'added' 'to'             : lambda: lambda x, y: x + y
    SUB -> 'minus'                           : lambda: lambda x, y: x - y
    MULT -> 'times' / 'multiplied' 'by'      : lambda: lambda x, y: x * y
    DIV -> 'divided' 'by'                    : lambda: lambda x, y: x / y
    """
)
```

Recall that in this grammar specification format, rules that are not explicitly provided with an augmentation (like all the OP rules after the first `OP -> ADD`) are associated with the textually most recent one (`lambda Op: Op`).

The `parse_augmented_grammar` function returns both an NLTK grammar and a dictionary that maps from productions in the grammar to their associated augmentations. Let's examine the returned grammar.

```
[ ]: for production in arithmetic_grammar.productions():
      print(f"{repr(production):25} {arithmetic_augmentations[production]}")
```

We can parse with the grammar using one of the built-in NLTK parsers.

```
[ ]: arithmetic_parser = nltk.parse.BottomUpChartParser(arithmetic_grammar)
      parses = [p for p in arithmetic_parser.parse('three plus one times four'.
      ↪split())]
      for parse in parses:
          parse.pretty_print()
```

Now let's turn to the augmentations. They can be arbitrary Python functions applied to the semantic representations associated with the right-hand-side nonterminals, returning the semantic representation of the left-hand side. To interpret the semantic representation of the entire sentence (at the root of the parse tree), we can use the following pseudo-code:

to interpret a tree:

```
    interpret each of the nonterminal-rooted subtrees
    find the augmentation associated with the root production of the tree
      (it should be a function of as many arguments as there are nonterminals on the right-hand side)
    return the result of applying the augmentation to the subtree values
```

(The base case of this recursion occurs when the number of nonterminal-rooted subtrees is zero, that is, a rule all of whose right-hand side elements are terminals.)

Suppose we had such a function, call it `interpret`. How would it operate on, for instance, the tree (S (S (NUM three)) (OP (ADD plus)) (S (NUM one)))?

```
interpret (S (S (NUM three)) (OP (ADD plus)) (S (NUM one)))
  |->interpret (S (NUM three))
  |   |->interpret (NUM three)
  |   |   |->(no subconstituents to evaluate)
  |   |   |->apply the augmentation for the rule NUM -> three to the empty set of values
  |   |   |   (lambda: 3) () ==> 3
  |   |   \==> 3
  |   |->apply the augmentation for the rule S -> NUM to the value 3
  |   |   (lambda NUM: NUM)(3) ==> 3
  |   \==> 3
  |->interpret (OP (ADD plus))
  |   |...
  |   \==> lambda x, y: x + y
  |->interpret (S (NUM one))
  |   |...
  |   \==> 1
  |->apply the augmentation for the rule S -> S OP S to the values 3, (lambda x, y: x + y), 1
  |   (lambda S1, Op, S2: Op(S1, S2))(3, (lambda x, y: x + y), 1) ==> 4
  \==> 4
```

Thus, the string “three plus one” is semantically interpreted as the value 4.

We provide the `interpret` function to carry out this recursive process, copied over from lab 4-2:

```
[ ]: def interpret(tree, augmentations):
      syntactic_rule = tree.productions()[0]
      semantic_rule = augmentations[syntactic_rule]
```

```

child_meanings = [interpret(child, augmentations)
                   for child in tree
                   if isinstance(child, nltk.Tree)]
return semantic_rule(*child_meanings)

```

Now we should be able to evaluate the arithmetic example from above.

```
[ ]: interpret(pauses[0], arithmetic_augmentations)
```

And we can even write a function that parses and interprets a string. We'll have it evaluate each of the possible parses and print the results.

```
[ ]: def parse_and_interpret(string, grammar, augmentations):
      parser = nltk.parse.BottomUpChartParser(grammar)
      parses = parser.parse(string.split())
      for parse in parses:
          parse.pretty_print()
          print(parse, "=>", interpret(parse, augmentations))

```

```
[ ]: parse_and_interpret("three plus one times four", arithmetic_grammar,
    ↪arithmetic_augmentations)
```

Since the string is syntactically ambiguous according to the grammar, it is semantically ambiguous as well.

## 2.2 Some grammar specification conveniences

Before going on, it will be useful to have a few more conveniences in writing augmentations for rules. First, since the augmentations are arbitrary Python expressions, they can be built from and make use of other functions. For instance, you'll notice that many of the augmentations at the leaves of the tree took no arguments and returned a constant. We can define a function `constant` that returns a function that ignores its arguments and returns a particular value.

```
[ ]: def constant(value):
      """Return `value`, ignoring any arguments"""
      return lambda *args: value

```

Similarly, several of the augmentations are functions that just return their first argument. Again, we can define a generic form `first` of such a function:

```
[ ]: def first(*args):
      """Return the value of the first (and perhaps only) subconstituent,
         ignoring any others"""
      return args[0]

```

We can now rewrite the grammar above to take advantage of these shortcuts.

In the call to `parse_augmented_grammar` below, we pass in the global environment, extracted via a `globals()` function call, via the named argument `globals`. This al-

lows the `parse_augmented_grammar` function to make use of the global bindings for `constant`, `first`, and the like when evaluating the augmentation expressions to their values. You can check out the code in `transform.py` to see how the passed in `globals` bindings are used. To help understand what's going on, see what happens if you don't include the `globals=globals()`.

```
[ ]: arithmetic_grammar_2, arithmetic_augmentations_2 = xform.
      ↪ parse_augmented_grammar(
          """
          ## Sample grammar for arithmetic expressions

          S -> NUM                                : first
              | S OP S                            : lambda S1, Op, S2: Op(S1, S2)

          OP -> ADD                                : first
              | SUB
              | MULT
              | DIV

          NUM -> 'zero'                            : constant(0)
              | 'one'                             : constant(1)
              | 'two'                             : constant(2)
              | 'three'                           : constant(3)
              | 'four'                            : constant(4)
              | 'five'                            : constant(5)
              | 'six'                             : constant(6)
              | 'seven'                           : constant(7)
              | 'eight'                           : constant(8)
              | 'nine'                            : constant(9)
              | 'ten'                             : constant(10)

          ADD -> 'plus' | 'added' 'to'             : constant(lambda x, y: x + y)
          SUB -> 'minus'                           : constant(lambda x, y: x - y)
          MULT -> 'times' | 'multiplied' 'by'      : constant(lambda x, y: x * y)
          DIV -> 'divided' 'by'                    : constant(lambda x, y: x / y)
          """,
          globals=globals())
```

Finally, it might make our lives easier to write a template of augmentations whose instantiation depends on the right-hand side of the rule.

We use a reserved keyword `_RHS` to denote the right-hand side of the syntactic rule, which will be replaced by a **list** of the right-hand-side strings. For example, an augmentation `numeric_template(_RHS)` would be as if written as `numeric_template(['zero'])` when the rule is `NUM -> 'zero'`, and `numeric_template(['one'])` when the rule is `NUM -> 'one'`. The details of how this works can be found at [scripts/transform.py](#).

This would allow us to use a single template function, for example,

```
[ ]: def numeric_template(rhs):
    """Ignore the subphrase meanings and lookup the first right-hand-side symbol
        as a number"""
    return constant({'zero':0, 'one':1, 'two':2, 'three':3, 'four':4, 'five':5,
                     'six':6, 'seven':7, 'eight':8, 'nine':9, 'ten':10}[rhs[0]])
```

and then further simplify the grammar specification:

```
[ ]: arithmetic_grammar_3, arithmetic_augmentations_3 = xform.
    ↪ parse_augmented_grammar(
        """
        ## Sample grammar for arithmetic expressions

        S -> NUM                                : first
          | S OP S                                : lambda S1, Op, S2: Op(S1, S2)

        OP -> ADD                                : first
          | SUB
          | MULT
          | DIV

        NUM -> 'zero' | 'one' | 'two'            : numeric_template(_RHS)
          | 'three' | 'four' | 'five'
          | 'six' | 'seven' | 'eight'
          | 'nine' | 'ten'

        ADD -> 'plus' | 'added' 'to'            : constant(lambda x, y: x + y)
        SUB -> 'minus'                          : constant(lambda x, y: x - y)
        MULT -> 'times' | 'multiplied' 'by'      : constant(lambda x, y: x * y)
        DIV -> 'divided' 'by'                    : constant(lambda x, y: x / y)
        """,
        globals=globals())

[ ]: parse_and_interpret("six divided by three", arithmetic_grammar_3,
    ↪ arithmetic_augmentations_3)
```

## 2.3 Example: *Green Eggs and Ham* revisited

This stuff is tricky, so it's useful to see more examples before jumping in the deep end. In this simple GEaH fragment grammar, we use a larger set of auxiliary functions to build the augmentations.

```
[ ]: def forward(F, A):
    """Forward application: Return the application of the first
        argument to the second"""
    return F(A)

def backward(A, F):
```



```

"""Backward application: Return the application of the second
    argument to the first"""
return F(A)

def second(*args):
    """Return the value of the second subconstituent, ignoring any others"""
    return args[1]

def ignore(*args):
    """Return `None`, ignoring everything about the constituent. (Good as a
        placeholder until a better augmentation can be devised.)"""
    return None

```

Using these, we can build and test the grammar.

```

[ ]: geah_grammar_spec = """
    ## Productions
    S -> NP VP          : backward
    VP -> V NP           : forward

    ## Lexicon
    V -> 'likes'         : constant(lambda Object: lambda Subject: λ
    ↪f"like({Subject}, {Object})")
    NP -> 'Sam' | 'sam'   : constant(_RHS[0])
    NP -> 'ham'
    NP -> 'eggs'
    """

```

```

[ ]: geah_grammar, geah_augmentations = xform.
    ↪parse_augmented_grammar(geah_grammar_spec,
    ↪globals=globals())

```

```

[ ]: parse_and_interpret("Sam likes ham", geah_grammar, geah_augmentations)

```

### 3 Semantics of ATIS queries

Now you're in a good position to understand and add augmentations to a more comprehensive grammar, say, one that parses ATIS queries and generates SQL queries.

In preparation for that, we need to load the ATIS data, both NL and SQL queries.

#### 3.1 Loading and preprocessing the corpus

To simplify things a bit, we'll only consider ATIS queries whose question type (remember that from project segment 1?) is `flight_id`. We download training, development, and test splits for this subset of the ATIS corpus, including corresponding SQL queries.

```
[ ]: # Acquire the datasets - training, development, and test splits of the
# ATIS queries and corresponding SQL queries
wget.download(f"{source_url}/ATIS/test_flightid.nl", out="data/")
wget.download(f"{source_url}/ATIS/test_flightid.sql", out="data/")
wget.download(f"{source_url}/ATIS/dev_flightid.nl", out="data/")
wget.download(f"{source_url}/ATIS/dev_flightid.sql", out="data/")
wget.download(f"{source_url}/ATIS/train_flightid.nl", out="data/")
wget.download(f"{source_url}/ATIS/train_flightid.sql", out="data/")
```

Let's take a look at the data: the NL queries are in .nl files, and the SQL queries are in .sql files.

```
[ ]: shell("head -1 data/dev_flightid.nl")
shell("head -1 data/dev_flightid.sql")
```

### 3.2 Corpus preprocessing

We'll use `torchtext` to process the data. We use two `Fields`: `SRC` for the questions, and `TGT` for the SQL queries. We'll use the tokenizer from project segment 3.

```
[ ]: ## Tokenizer
tokenizer = nltk.tokenize.RegexpTokenizer('\d+|st\.|[\w-]+\.[\d\.]+\S+')
def tokenize(string):
    return tokenizer.tokenize(string.lower())

## Demonstrating the tokenizer
## Note especially the handling of `11pm` and hyphenated words.
print(tokenize("Are there any first-class flights from St. Louis at 11pm for_
→less than $3.50?"))

[ ]: SRC = tt.data.Field(include_lengths=True,           # include lengths
                        batch_first=False,             # batches will be max_len x_
                        →batch_size
                        tokenize=tokenize,             # use our tokenizer
                        )
TGT = tt.data.Field(include_lengths=False,             # batches will be max_len x_
                    →batch_size
                    tokenize=lambda x: x.split(),     # use split to tokenize
                    init_token="<bos>",               # prepend <bos>
                    eos_token="<eos>",               # append <eos>
                    fields = [('src', SRC), ('tgt', TGT)]
```

Note that we specified `batch_first=False` (as in lab 4-4), so that the returned batched tensors would be of size `max_length x batch_size`, which facilitates `seq2seq` implementation.

Now, we load the data using `torchtext`. We use the `TranslationDataset` class here because our task is essentially a translation task: “translating” questions into the corresponding SQL queries.

Therefore, we also refer to the questions as the *source* side (SRC) and the SQL queries as the *target* side (TGT).

```
[ ]: # Make splits for data
train_data, val_data, test_data = tt.datasets.TranslationDataset.splits(
    ('_flightid.nl', '_flightid.sql'), fields, path='./data/',
    train='train', validation='dev', test='test')

MIN_FREQ = 3
SRC.build_vocab(train_data.src, min_freq=MIN_FREQ)
TGT.build_vocab(train_data.tgt, min_freq=MIN_FREQ)

print (f"Size of English vocab: {len(SRC.vocab)}")
print (f"Most common English words: {SRC.vocab.freqs.most_common(10)}\n")

print (f"Size of SQL vocab: {len(TGT.vocab)}")
print (f"Most common SQL words: {TGT.vocab.freqs.most_common(10)}\n")

print (f"Index for start of sequence token: {TGT.vocab.stoi[TGT.init_token]}")
print (f"Index for end of sequence token: {TGT.vocab.stoi[TGT.eos_token]}")
```

Next, we batch our data to facilitate processing on a GPU. Batching is a bit tricky because the source and target will typically be of different lengths. Fortunately, `torchtext` allows us to pass in a `sort_key` function. By sorting on length, we can minimize the amount of padding on the source side, but since there is still some padding, we need to handle them with `pack` and `unpack` later on in the `seq2seq` part (as in lab 4-5).

```
[ ]: BATCH_SIZE = 16 # batch size for training/validation
TEST_BATCH_SIZE = 1 # batch size for test, we use 1 to make beam search
    ↪ implementation easier

train_iter, val_iter = tt.data.BucketIterator.splits((train_data, val_data),
                                                    batch_size=BATCH_SIZE,
                                                    device=device,
                                                    repeat=False,
                                                    sort_key=lambda x: len(x.
    ↪src),
                                                    sort_within_batch=True)

test_iter = tt.data.BucketIterator(test_data,
                                   batch_size=TEST_BATCH_SIZE,
                                   device=device,
                                   repeat=False,
                                   sort=False,
                                   train=False)
```

Let's look at a single batch from one of these iterators.

```
[ ]: batch = next(iter(train_iter))
train_batch_text, train_batch_text_lengths = batch.src
print (f"Size of text batch: {train_batch_text.shape}")
print (f"Third sentence in batch: {train_batch_text[:, 2]}")
print (f"Length of the third sentence in batch: {train_batch_text_lengths[2]}")
print (f"Converted back to string: {' '.join([SRC.vocab.itos[i] for i in
→train_batch_text[:, 2]])}")

train_batch_sql = batch.tgt
print (f"Size of sql batch: {train_batch_sql.shape}")
print (f"Third SQL in batch: {train_batch_sql[:, 2]}")
print (f"Converted back to string: {' '.join([TGT.vocab.itos[i] for i in
→train_batch_sql[:, 2]])}")
```

Alternatively, we can directly iterate over the raw examples:

```
[ ]: for example in train_iter.dataset[:1]:
    train_text_1 = ' '.join(example.src) # detokenized question
    train_sql_1 = ' '.join(example.tgt)  # detokenized sql
    print (f"Question: {train_text_1}\n")
    print (f"SQL: {train_sql_1}")
```

### 3.3 Establishing a SQL database for evaluating ATIS queries

The output of our systems will be SQL queries. How should we determine if the generated queries are correct? We can't merely compare against the gold SQL queries, since there are many ways to implement a SQL query that answers any given NL query.

Instead, we will execute the queries – both the predicted SQL query and the gold SQL query – on an actual database, and verify that the returned responses are the same. For that purpose, we need a SQL database server to use. We'll set one up here, using the [Python sqlite3 module](#).

```
[ ]: @func_set_timeout(TIMEOUT)
def execute_sql(sql):
    conn = sqlite3.connect('data/atis_sqlite.db') # establish the DB based on
→the downloaded data
    c = conn.cursor() # build a "cursor"
    c.execute(sql)
    results = list(c.fetchall())
    c.close()
    conn.close()
    return results
```

To run a query, we use the cursor's `execute` function, and retrieve the results with `fetchall`. Let's get all the flights that arrive at General Mitchell International – the query `train_sql_1` above. There's a lot, so we'll just print out the first few.

```
[ ]: predicted_ret = execute_sql(train_sql_1)

print(f"""
Executing: {train_sql_1}

Result: {len(predicted_ret)} entries starting with

{predicted_ret[:10]}
""")
```

For your reference, the SQL database we are using has a database schema described at <https://github.com/jkkummerfeld/text2sql-data/blob/master/data/atis-schema.csv>, and is consistent with the SQL queries provided in the various `.sql` files loaded above.

## 4 Rule-based parsing and interpretation of ATIS queries

First, you will implement a rule-based semantic parser using a grammar like the one you completed in the third project segment. We've placed an initial grammar in the file `data/grammar`. In addition to the helper functions defined above (`constant`, `first`, etc.), it makes use of some other simple functions. We've included those below, but you can (and almost certainly should) augment this set with others that you define as you build out the full set of augmentations.

```
[ ]: def upper(term):
    return '' + term.upper() + ''

def weekday(day):
    return f"flight.flight_days IN (SELECT days.days_code FROM days WHERE days.
    ↳day_name = '{day.upper()}')"

def month_name(month):
    return {'JANUARY' : 1,
            'FEBRUARY' : 2,
            'MARCH' : 3,
            'APRIL' : 4,
            'MAY' : 5,
            'JUNE' : 6,
            'JULY' : 7,
            'AUGUST' : 8,
            'SEPTEMBER' : 9,
            'OCTOBER' : 10,
            'NOVEMBER' : 11,
            'DECEMBER' : 12}[month.upper()]

def airports_from_airport_name(airport_name):
    return f"(SELECT airport.airport_code FROM airport WHERE airport.airport_name_
    ↳= {upper(airport_name)})"
```

```

def airports_from_city(city):
    return f"""
        (SELECT airport_service.airport_code FROM airport_service WHERE
        ↪airport_service.city_code IN
            (SELECT city.city_code FROM city WHERE city.city_name = {upper(city)}))
        """

def null_condition(*args, **kwargs):
    return 1

def depart_around(time):
    return f"""
        flight.departure_time >= {add_delta(miltime(time), -15).strftime('%H%M')}
        AND flight.departure_time <= {add_delta(miltime(time), 15).strftime('%H%M')}
        """.strip()

def add_delta(tme, delta):
    # transform to a full datetime first
    return (datetime.datetime.combine(datetime.date.today(), tme) +
            datetime.timedelta(minutes=delta)).time()

def miltime(minutes):
    return datetime.time(hour=int(minutes/100), minute=(minutes % 100))

```

We can build a parser with the augmented grammar:

```

[ ]: atis_grammar, atis_augmentations = xform.read_augmented_grammar('data/grammar',
    ↪globals=globals())
    atis_parser = nltk.parse.BottomUpChartParser(atis_grammar)

```

We'll define a function to return a parse tree for a string according to the ATIS grammar (if available).

```

[ ]: def parse_tree(sentence):
    """Parse a sentence and return the parse tree, or None if failure."""
    try:
        parses = list(atis_parser.parse(tokenize(sentence)))
        if len(parses) == 0:
            return None
        else:
            return parses[0]
    except:
        return None

```

We can check the overall coverage of this grammar on the training set by using the `parse_tree` function to determine if a parse is available. The grammar that we provide should get about a 40% coverage of the training set.

```
[ ]: # Check coverage on training set
parsed = 0
with open("data/train_flightid.nl") as train:
    examples = train.readlines()[:]
    for sentence in tqdm(examples):
        if parse_tree(sentence):
            parsed += 1
        else:
            next

print(f"\nParsed {parsed} of {len(examples)} ({parsed*100/(len(examples)):.
    ↳2f}%")
```

#### 4.1 Goal 1: Construct SQL queries from a parse tree and evaluate the results

It's time to turn to the first major part of this project segment, implementing a rule-based semantic parsing system to answer flight-ID-type ATIS queries.

Recall that in rule-based semantic parsing, each syntactic rule is associated with a semantic composition rule. The grammar we've provided has semantic augmentations for some of the low-level phrases – cities, airports, times, airlines – but not the higher level syntactic types. You'll be adding those.

In the ATIS grammar that we provide, as with the earlier toy grammars, the augmentation for a rule with  $n$  nonterminals and  $m$  terminals on the right-hand side is assumed to be called with  $n$  positional arguments (the values for the corresponding children). The `interpret` function you've already defined should therefore work well with this grammar.

Let's run through one way that a semantic derivation might proceed, for the sample query "flights to boston":

```
[ ]: sample_query = "flights to boston"
print(tokenize(sample_query))
sample_tree = parse_tree(sample_query)
sample_tree.pretty_print()
```

Given a sentence, we first construct its parse tree using the syntactic rules, then compose the corresponding semantic rules bottom-up, until eventually we arrive at the root node with a finished SQL statement. For this query, we will go through what the possible meaning representations for the constituents of "flights to boston" might be. But this is just one way of doing things; other ways are possible, and you should feel free to experiment.

Working from bottom up:

1. The `TERM_PLACE` phrase "boston" uses the composition function template `constant(airports_from_city(' '.join(_RHS)))`, which will be instantiated as `constant(airports_from_city(' '.join(['boston'])))` (recall that `_RHS` is replaced by the right-hand side of the rule). The meaning of `TERM_PLACE` will be the SQL snippet

```
SELECT airport_service.airport_code
FROM airport_service
```

```
WHERE airport_service.city_code IN
  (SELECT city.city_code
   FROM city
   WHERE city.city_name = "BOSTON")
```

(This query generates a list of all of the airports in Boston.)

2. The N\_PLACE phrase “boston” can have the same meaning as the TERM\_PLACE.
3. The P\_PLACE phrase “to” might be associated with a function that maps a SQL query for a list of airports to a SQL condition that holds of flights that go to one of those airports, i.e., `flight.to_airport IN (...)`.
4. The PP\_PLACE phrase “to boston” might apply the P\_PLACE meaning to the TERM\_PLACE meaning, thus generating a SQL condition that holds of flights that go to one of the Boston airports:

```
flight.to_airport IN
  (SELECT airport_service.airport_code
   FROM airport_service
   WHERE airport_service.city_code IN
     (SELECT city.city_code
      FROM city
      WHERE city.city_name = "BOSTON"))
```

5. The PP phrase “to Boston” can again get its meaning from the PP\_PLACE.
6. The TERM\_FLIGHT phrase “flights” might also return a condition on flights, this time the “null condition”, represented by the SQL truth value 1. Ditto for the N\_FLIGHT phrase “flights”.
7. The N\_FLIGHT phrase “flights to boston” can conjoin the two conditions, yielding the SQL condition

```
flight.to_airport IN
  (SELECT airport_service.airport_code
   FROM airport_service
   WHERE airport_service.city_code IN
     (SELECT city.city_code
      FROM city
      WHERE city.city_name = "BOSTON"))
```

AND 1

which can be inherited by the NOM\_FLIGHT and NP\_FLIGHT phrases.

8. The S phrase “flights to boston” can use the condition provided by the NP\_FLIGHT phrase to select all flights satisfying the condition with a SQL query like

```
SELECT DISTINCT flight.flight_id
FROM flight
WHERE flight.to_airport IN
  (SELECT airport_service.airport_code
   FROM airport_service
   WHERE airport_service.city_code IN
     (SELECT city.city_code
```



```

        FROM city
        WHERE city.city_name = "BOSTON")
    AND 1

```

This SQL query is then taken to be a representation of the meaning for the NL query “flights to boston”, and can be executed against the ATIS database to retrieve the requested flights.

Now, it’s your turn to add augmentations to `data/grammar` to make this example work. The augmentations that we have provided for the grammar make use of a set of auxiliary functions that we defined above. You should feel free to add your own auxiliary functions that you make use of in the grammar.

```

[ ]: #TODO: add augmentations to `data/grammar` to make this example work
    atis_grammar, atis_augmentations = xform.read_augmented_grammar('data/grammar',
↪globals=globals())
    atis_parser = nltk.parse.BottomUpChartParser(atis_grammar)
    predicted_sql = interpret(sample_tree, atis_augmentations)
    print("Predicted SQL:\n\n", predicted_sql, "\n")

```

**Verification on some examples** With a rule-based semantic parsing system, we can generate SQL queries given questions, and then execute those queries on a SQL database to answer the given questions. To evaluate the performance of the system, we compare the returned results against the results of executing the ground truth queries.

We provide a function `verify` to compare the results from our generated SQL to the ground truth SQL. It should be useful for testing individual queries.

```

[ ]: def verify(predicted_sql, gold_sql, silent=True):
    """
    Compare the correctness of the generated SQL by executing on the
    ATIS database and comparing the returned results.
    Arguments:
        predicted_sql: the predicted SQL query
        gold_sql: the reference SQL query to compare against
        silent: print outputs or not
    Returns: True if the returned results are the same, otherwise False
    """
    # Execute predicted SQL
    try:
        predicted_result = execute_sql(predicted_sql)
    except BaseException as e:
        if not silent:
            print(f"predicted sql exec failed: {e}")
        return False
    if not silent:
        print("Predicted DB result:\n\n", predicted_result[:10], "\n")

    # Execute gold SQL
    try:

```

```

    gold_result = execute_sql(gold_sql)
except BaseException as e:
    if not silent:
        print(f"gold sql exec failed: {e}")
    return False
if not silent:
    print("Gold DB result:\n\n", gold_result[:10], "\n")

# Verify correctness
if gold_result == predicted_result:
    return True

```

Let's try this methodology on a simple example: "flights from phoenix to milwaukee". we provide it along with the gold SQL query.

```

[ ]: def rule_based_trial(sentence, gold_sql):
    print("Sentence: ", sentence, "\n")
    tree = parse_tree(sentence)
    print("Parse:\n\n")
    tree.pretty_print()

    predicted_sql = interpret(tree, atis_augmentations)
    print("Predicted SQL:\n\n", predicted_sql, "\n")

    if verify(predicted_sql, gold_sql, silent=False):
        print ('Correct!')
    else:
        print ('Incorrect!')

```

```

[ ]: # Run this cell to reload augmentations after you make changes to `data/grammar`
atis_grammar, atis_augmentations = xform.read_augmented_grammar('data/grammar',
↪globals=globals())
atis_parser = nltk.parse.BottomUpChartParser(atis_grammar)

```

```

[ ]: #TODO: add augmentations to `data/grammar` to make this example work
# Example 1
example_1 = 'flights from phoenix to milwaukee'
gold_sql_1 = """
SELECT DISTINCT flight_1.flight_id
FROM flight flight_1 ,
     airport_service airport_service_1 ,
     city city_1 ,
     airport_service airport_service_2 ,
     city city_2
WHERE flight_1.from_airport = airport_service_1.airport_code
     AND airport_service_1.city_code = city_1.city_code
     AND city_1.city_name = 'PHOENIX'

```

```

        AND flight_1.to_airport = airport_service_2.airport_code
        AND airport_service_2.city_code = city_2.city_code
        AND city_2.city_name = 'MILWAUKEE'
    """

```

```
rule_based_trial(example_1, gold_sql_1)
```

To make development faster, we recommend starting with a few examples before running the full evaluation script. We've taken some examples from the ATIS dataset including the gold SQL queries that they provided. Of course, yours (and those of the project segment solution set) may differ.

```

[ ]: #TODO: add augmentations to `data/grammar` to make this example work
# Example 2
example_2 = 'i would like a united flight'
gold_sql_2 = """
    SELECT DISTINCT flight_1.flight_id
    FROM flight flight_1
    WHERE flight_1.airline_code = 'UA'
    """

rule_based_trial(example_2, gold_sql_2)

```

```

[ ]: #TODO: add augmentations to `data/grammar` to make this example work
# Example 3
example_3 = 'i would like a flight between boston and dallas'
gold_sql_3 = """
    SELECT DISTINCT flight_1.flight_id
    FROM flight flight_1 ,
        airport_service airport_service_1 ,
        city city_1 ,
        airport_service airport_service_2 ,
        city city_2
    WHERE flight_1.from_airport = airport_service_1.airport_code
        AND airport_service_1.city_code = city_1.city_code
        AND city_1.city_name = 'BOSTON'
        AND flight_1.to_airport = airport_service_2.airport_code
        AND airport_service_2.city_code = city_2.city_code
        AND city_2.city_name = 'DALLAS'
    """

# Note that the parse tree might appear wrong: instead of
# `PP_PLACE -> 'between' N_PLACE 'and' N_PLACE`, the tree appears to be
# `PP_PLACE -> 'between' 'and' N_PLACE N_PLACE`. But it's only a visualization
# error of tree.pretty_print() and you should assume that the production is
# `PP_PLACE -> 'between' N_PLACE 'and' N_PLACE` (you can verify by printing out
# all productions).

```

```
rule_based_trial(example_3, gold_sql_3)
```

```
[ ]: #TODO: add augmentations to `data/grammar` to make this example work
# Example 4
example_4 = 'show me the united flights from denver to baltimore'
gold_sql_4 = """
SELECT DISTINCT flight_1.flight_id
FROM flight flight_1 ,
     airport_service airport_service_1 ,
     city city_1 ,
     airport_service airport_service_2 ,
     city city_2
WHERE flight_1.airline_code = 'UA'
     AND ( flight_1.from_airport = airport_service_1.airport_code
           AND airport_service_1.city_code = city_1.city_code
           AND city_1.city_name = 'DENVER'
           AND flight_1.to_airport = airport_service_2.airport_code
           AND airport_service_2.city_code = city_2.city_code
           AND city_2.city_name = 'BALTIMORE' )

"""

rule_based_trial(example_4, gold_sql_4)
```

```
[ ]: #TODO: add augmentations to `data/grammar` to make this example work
# Example 5
example_5 = 'show flights from cleveland to miami that arrive before 4pm'
gold_sql_5 = """
SELECT DISTINCT flight_1.flight_id
FROM flight flight_1 ,
     airport_service airport_service_1 ,
     city city_1 ,
     airport_service airport_service_2 ,
     city city_2
WHERE flight_1.from_airport = airport_service_1.airport_code
     AND airport_service_1.city_code = city_1.city_code
     AND city_1.city_name = 'CLEVELAND'
     AND ( flight_1.to_airport = airport_service_2.airport_code
           AND airport_service_2.city_code = city_2.city_code
           AND city_2.city_name = 'MIAMI'
           AND flight_1.arrival_time < 1600 )

"""

rule_based_trial(example_5, gold_sql_5)
```

```
[ ]: #TODO: add augmentations to `data/grammar` to make this example work
# Example 6
```

```

example_6 = 'okay how about a flight on sunday from tampa to charlotte'
gold_sql_6 = """
SELECT DISTINCT flight_1.flight_id
FROM flight flight_1 ,
      airport_service airport_service_1 ,
      city city_1 ,
      airport_service airport_service_2 ,
      city city_2 ,
      days days_1 ,
      date_day date_day_1
WHERE flight_1.from_airport = airport_service_1.airport_code
      AND airport_service_1.city_code = city_1.city_code
      AND city_1.city_name = 'TAMPA'
      AND ( flight_1.to_airport = airport_service_2.airport_code
            AND airport_service_2.city_code = city_2.city_code
            AND city_2.city_name = 'CHARLOTTE'
            AND flight_1.flight_days = days_1.days_code
            AND days_1.day_name = date_day_1.day_name
            AND date_day_1.year = 1991
            AND date_day_1.month_number = 8
            AND date_day_1.day_number = 27 )

"""

# You might notice that the gold answer above used the exact date, which is
# not easily implementable. A more implementable way (generated by the project
# segment 4 solution code) is:
gold_sql_6b = """
SELECT DISTINCT flight.flight_id
FROM flight
WHERE (((1
      AND flight.flight_days IN (SELECT days.days_code
                                FROM days
                                WHERE days.day_name = 'SUNDAY')
      )
      AND flight.from_airport IN (SELECT airport_service.airport_code
                                FROM airport_service
                                WHERE airport_service.city_code IN_
↪(SELECT city.city_code
                                FROM_
↪city
                                _
↪WHERE city.city_name = "TAMPA"))))
      AND flight.to_airport IN (SELECT airport_service.airport_code
                                FROM airport_service
                                WHERE airport_service.city_code IN (SELECT_
↪city.city_code

```

```

                                FROM_
↪city

                                WHERE_
↪city.city_name = "CHARLOTTE"))))
    ""

rule_based_trial(example_6, gold_sql_6b)

```

```

[ ]: #TODO: add augmentations to `data/grammar` to make this example work
# Example 7
example_7 = 'list all flights going from boston to atlanta that leaves before 7_
↪am on thursday'
gold_sql_7 = """
SELECT DISTINCT flight_1.flight_id
FROM flight flight_1 ,
     airport_service airport_service_1 ,
     city city_1 ,
     airport_service airport_service_2 ,
     city city_2 ,
     days days_1 ,
     date_day date_day_1
WHERE flight_1.from_airport = airport_service_1.airport_code
     AND airport_service_1.city_code = city_1.city_code
     AND city_1.city_name = 'BOSTON'
     AND ( flight_1.to_airport = airport_service_2.airport_code
           AND airport_service_2.city_code = city_2.city_code
           AND city_2.city_name = 'ATLANTA'
           AND ( flight_1.flight_days = days_1.days_code
                 AND days_1.day_name = date_day_1.day_name
                 AND date_day_1.year = 1991
                 AND date_day_1.month_number = 5
                 AND date_day_1.day_number = 24
                 AND flight_1.departure_time < 700 ) )

""

# Again, the gold answer above used the exact date, as opposed to the
# following approach:
gold_sql_7b = """
SELECT DISTINCT flight.flight_id
FROM flight
WHERE ((1
      AND (((1
            AND flight.from_airport IN (SELECT airport_service.
↪airport_code

                                FROM airport_service
                                WHERE airport_service.city_code_
↪IN (SELECT city.city_code

```

```

↪ FROM city

↪ WHERE city.city_name = "BOSTON"))
        AND flight.to_airport IN (SELECT airport_service.airport_code
                                FROM airport_service
                                WHERE airport_service.city_code IN_
↪(SELECT city.city_code
                                FROM city

↪WHERE city.city_name = "ATLANTA"))
        AND flight.departure_time <= 0700)
        AND flight.flight_days IN (SELECT days.days_code
                                FROM days
                                WHERE days.day_name = 'THURSDAY'))))

"""

rule_based_trial(example_7, gold_sql_7b)

```

```

[ ]: #TODO: add augmentations to `data/grammar` to make this example work
# Example 8
example_8 = 'list the flights from dallas to san francisco on american airlines'
gold_sql_8 = """
SELECT DISTINCT flight_1.flight_id
FROM flight flight_1 ,
     airport_service airport_service_1 ,
     city city_1 ,
     airport_service airport_service_2 ,
     city city_2
WHERE flight_1.airline_code = 'AA'
     AND ( flight_1.from_airport = airport_service_1.airport_code
         AND airport_service_1.city_code = city_1.city_code
         AND city_1.city_name = 'DALLAS'
         AND flight_1.to_airport = airport_service_2.airport_code
         AND airport_service_2.city_code = city_2.city_code
         AND city_2.city_name = 'SAN FRANCISCO' )

"""

rule_based_trial(example_8, gold_sql_8)

```

#### 4.1.1 Systematic evaluation on a test set

We can perform a more systematic evaluation by checking the accuracy of the queries on an entire test set for which we have gold queries. The `evaluate` function below does just this, calculating precision, recall, and F1 metrics for the test set. It takes as argument a “predictor” function, which maps token sequences to predicted SQL queries. We’ve provided a predictor function for the

rule-based model in the next cell (and a predictor for the seq2seq system below when we get to that system).

The rule-based system does not generate predictions for all queries; many queries won't parse. The precision and recall metrics take this into account in measuring the efficacy of the method. The recall metric captures what proportion of *all of the test examples* for which the system generates a correct query. The precision metric captures what proportion of *all of the test examples for which a prediction is generated* for which the system generates a correct query. (Recall that F1 is just the geometric mean of precision and recall.)

Once you've made some progress on adding augmentations to the grammar, you can evaluate your progress by seeing if the precision and recall have improved. For reference, the solution code achieves precision of about 66% and recall of about 28% for an F1 of 39%.

```
[ ]: def evaluate(predictor, dataset, num_examples=0, silent=True):
    """Evaluate accuracy of `predictor` by executing predictions on a
    SQL database and comparing returned results against those of gold queries.

    Arguments:
        predictor: a function that maps a token sequence (provided by
        → torchtext)
                   to a predicted SQL query string
        dataset: the dataset of token sequences and gold SQL queries
        num_examples: number of examples from `dataset` to use; all of
                      them if 0
        silent: if set to False, will print out logs
    Returns: precision, recall, and F1 score
    """
    # Prepare to count results
    if num_examples <= 0:
        num_examples = len(dataset)
    example_count = 0
    predicted_count = 0
    correct = 0
    incorrect = 0

    # Process the examples from the dataset
    for example in tqdm(dataset[:num_examples]):
        example_count += 1
        # obtain query SQL
        predicted_sql = predictor(example.src)
        if predicted_sql == None:
            continue
        predicted_count += 1
        # obtain gold SQL
        gold_sql = ' '.join(example.tgt)

        # check that they're compatible
```



```

    if verify(predicted_sql, gold_sql):
        correct += 1
    else:
        incorrect += 1

    # Compute and return precision, recall, F1
    precision = correct / predicted_count if predicted_count > 0 else 0
    recall = correct / example_count
    f1 = (2 * precision * recall) / (precision + recall) if precision + recall > 0
    ↪ else 0
    return precision, recall, f1

```

```

[ ]: def rule_based_predictor(tokens):
    query = ' '.join(tokens)    # detokenized query
    tree = parse_tree(query)
    if tree is None:
        return None
    try:
        predicted_sql = interpret(tree, atis_augmentations)
    except Exception as err:
        return None
    return predicted_sql

```

```

[ ]: precision, recall, f1 = evaluate(rule_based_predictor, test_iter.dataset,
    ↪ num_examples=0)
print(f"precision: {precision:3.2f}")
print(f"recall:    {recall:3.2f}")
print(f"F1:       {f1:3.2f}")

```

## 5 End-to-End Seq2Seq Model

In this part, you will implement a seq2seq model **with attention mechanism** to directly learn the translation from NL query to SQL. You might find labs 4-4 and 4-5 particularly helpful, as the primary difference here is that we are using a different dataset.

**Note:** We recommend using GPUs to train the model in this part (one way to get GPUs is to use [Google Colab](#) and clicking Menu -> Runtime -> Change runtime type -> GPU), as we need to use a very large model to solve the task well. For development we recommend starting with a smaller model and training for only 1 epoch.

### 5.1 Goal 2: Implement a seq2seq model (with attention)

In lab 4-5, you implemented a neural encoder-decoder model with attention. That model was used to convert English number phrases to numbers, but one of the biggest advantages of neural models is that we can easily apply them to different tasks (such as machine translation and document summarization) by using different training datasets.

Implement the class `AttnEncoderDecoder` to convert natural language queries into SQL statements.

You may find that you can reuse most of the code you wrote for lab 4-5. A reasonable way to proceed is to implement the following methods:

- **Model**

1. `__init__`: an initializer where you create network modules.
2. `forward`: given source word ids of size `(max_src_len, batch_size)`, source lengths of size `(batch_size)` and decoder input target word ids `(max_tgt_len, batch_size)`, returns logits `(max_tgt_len, batch_size, V_tgt)`. For better modularity you might want to implement it by implementing two functions `forward_encoder` and `forward_decoder`.

- **Optimization**

3. `train_all`: compute loss on training data, compute gradients, and update model parameters to minimize the loss.
4. `evaluate_ppl`: evaluate the current model's perplexity on a given dataset iterator, we use the perplexity value on the validation set to select the best model.

- **Decoding**

5. `predict`: Generates the target sequence given a list of source tokens using beam search decoding. Note that here you can assume the batch size to be 1 for simplicity.

```
[ ]: #TODO - implement the `AttnEncoderDecoder` class.
class AttnEncoderDecoder(nn.Module):
    ...
```

```
[ ]: ...
```

We provide the recommended hyperparameters for the final model in the script below, but you are free to tune the hyperparameters or change any part of the provided code.

For quick debugging, we recommend starting with smaller models (by using a very small `hidden_size`), and only a single epoch. If the model runs smoothly, then you can train the full model on GPUs.

```
[ ]: EPOCHS = 50 # epochs; we recommend starting with a smaller number like 1
LEARNING_RATE = 1e-4 # learning rate

# Instantiate and train classifier
model = AttnEncoderDecoder(SRC, TGT,
    hidden_size    = 1024,
    layers         = 1,
).to(device)

model.train_all(train_iter, val_iter, epochs=EPOCHS,
    ↪learning_rate=LEARNING_RATE)
model.load_state_dict(model.best_model)
```

```
# Evaluate model performance, the expected value should be < 1.2
print (f'Validation perplexity: {model.evaluate_ppl(val_iter):.3f}')
```

With a trained model, we can convert questions to SQL statements. We recommend making sure that the model can generate at least reasonable results on the examples from before, before evaluating on the full test set.

```
[ ]: def seq2seq_trial(sentence, gold_sql):
    print("Sentence: ", sentence, "\n")
    tokens = tokenize(sentence)

    predicted_sql = model.predict(tokens, K=1, max_T=400)
    print("Predicted SQL:\n\n", predicted_sql, "\n")

    if verify(predicted_sql, gold_sql, silent=False):
        print ('Correct!')
    else:
        print ('Incorrect!')
```

```
[ ]: seq2seq_trial(example_1, gold_sql_1)
```

```
[ ]: seq2seq_trial(example_2, gold_sql_2)
```

```
[ ]: seq2seq_trial(example_3, gold_sql_3)
```

```
[ ]: seq2seq_trial(example_4, gold_sql_4)
```

```
[ ]: seq2seq_trial(example_5, gold_sql_5)
```

```
[ ]: seq2seq_trial(example_6, gold_sql_6b)
```

```
[ ]: seq2seq_trial(example_7, gold_sql_7b)
```

```
[ ]: seq2seq_trial(example_8, gold_sql_8)
```

### 5.1.1 Evaluation

Now we are ready to run the full evaluation. A proper implementation should reach more than 35% precision/recall/F1.

```
[ ]: def seq2seq_predictor(tokens):
    prediction = model.predict(tokens, K=1, max_T=400)
    return prediction
```

```
[ ]: precision, recall, f1 = evaluate(seq2seq_predictor, test_iter.dataset,
    ↪ num_examples=0)
print(f"precision: {precision:3.2f}")
```

```
print(f"recall:    {recall:3.2f}")
print(f"F1:       {f1:3.2f}")
```

## 5.2 Goal 3: Implement a seq2seq model (with cross attention and self attention)

In the previous section, you have implemented a seq2seq model with attention. The attention mechanism used in that section is usually referred to as “cross-attention”, as at each decoding step, the decoder attends to encoder outputs, enabling a dynamic view on the encoder side as decoding proceeds.

Similarly, we can have a dynamic view on the decoder side as well as decoding proceeds, i.e., the decoder attends to decoder outputs at previous steps. This is called “self attention”, and has been found very useful in modern neural architectures such as transformers.

Augment the seq2seq model you implemented before with a decoder self-attention mechanism as class `AttnEncoderDecoder2`. A model diagram can be found below:

At each decoding step, the decoder LSTM first produces an output state  $o_t$ , then it attends to all previous output states  $o_1, \dots, o_{t-1}$  (decoder self-attention). You need to special case the first decoding step to not perform self-attention, as there are no previous decoder states. The attention result is added to  $o_t$  itself and the sum is used as  $q_t$  to attend to the encoder side (encoder-decoder cross-attention). The rest of the model is the same as encoder-decoder with attention.

```
[ ]: #TODO - implement the `AttnEncoderDecoder2` class.
class AttnEncoderDecoder2(nn.Module):
    ...
```

```
[ ]: ...
```

```
[ ]: EPOCHS = 50 # epochs, we recommend starting with a smaller number like 1
LEARNING_RATE = 1e-4 # learning rate

# Instantiate and train classifier
model2 = AttnEncoderDecoder2(SRC, TGT,
    hidden_size    = 1024,
    layers         = 1,
).to(device)

model2.train_all(train_iter, val_iter, epochs=EPOCHS,
    ↪learning_rate=LEARNING_RATE)
model2.load_state_dict(model2.best_model)

# Evaluate model performance, the expected value should be < 1.2
print (f'Validation perplexity: {model2.evaluate_ppl(val_iter):.3f}')
```

### 5.2.1 Evaluation

Now we are ready to run the full evaluation. A proper implementation should reach more than 35% precision/recall/F1.

```
[ ]: def seq2seq_predictor2(tokens):
      prediction = model2.predict(tokens, K=1, max_T=400)
      return prediction

[ ]: precision, recall, f1 = evaluate(seq2seq_predictor2, test_iter.dataset,
      ↪num_examples=0)
print(f"precision: {precision:3.2f}")
print(f"recall:    {recall:3.2f}")
print(f"F1:       {f1:3.2f}")
```

## 6 Discussion

### 6.1 Goal 4: Compare the pros and cons of rule-based and neural approaches.

Compare the pros and cons of the rule-based approach and the neural approaches with relevant examples from your experiments above. Concerning the accuracy, which approach would you choose to be used in a product? Explain.

*Type your answer here, replacing this text.*

### 6.2 (Optional) Goal 5: Use state-of-the-art pretrained transformers

The most recent breakthrough in natural-language processing stems from the use of pretrained transformer models. For example, you might have heard of pretrained transformers such as [GPT-3](#) and [BERT](#). (BERT is already used in [Google search](#).) These models are usually trained on vast amounts of text data using variants of language modeling objectives, and researchers have found that finetuning them on downstream tasks usually results in better performance as compared to training a model from scratch.

In the previous part, you implemented an LSTM-based sequence-to-sequence approach. To “upgrade” the model to be a state-of-the-art pretrained transformer only requires minor modifications.

The pretrained model that we will use is [BART](#), which uses a bidirectional transformer encoder and a unidirectional transformer decoder, as illustrated in the below diagram (image courtesy <https://arxiv.org/pdf/1910.13461>):

We can see that this model is strikingly similar to the LSTM-based encoder-decoder model we’ve been using. The only difference is that they use transformers instead of LSTMs. Therefore, we only need to change the modeling parts of the code, as we will see later.

First, we download and load the pretrained BART model from the [transformers](#) package by Huggingface. Note that we also need to use the “tokenizer” of BART, which is actually a combination of a tokenizer and a mapping from strings to word ids.

```
[ ]: pretrained_bart = BartForConditionalGeneration.from_pretrained('facebook/
      ↪bart-base')
bart_tokenizer = BartTokenizer.from_pretrained('facebook/bart-base')
```

Below we demonstrate how to use BART’s tokenizer to convert a sentence to a list of word ids, and vice versa.

```
[ ]: # BART uses a predefined "tokenizer", which directly maps a sentence
# to a list of ids
def bart_tokenize(string):
    return bart_tokenizer(string)['input_ids'][:1024] # BART model can process at
    ↳ most 1024 tokens

def bart_detokenize(token_ids):
    return bart_tokenizer.decode(token_ids, skip_special_tokens=True)

## Demonstrating the tokenizer
question = 'Are there any first-class flights from St. Louis at 11pm for less
    ↳ than $3.50?'

tokenized_question = bart_tokenize(question)
print('tokenized:', tokenized_question)

detokenized_question = bart_detokenize(tokenized_question)
print('detokenized:', detokenized_question)
```

We need to reprocess the data using our new tokenizer. Note that here we set `batch_first` to `True`, since that's the expected input shape of the transformers package.

```
[ ]: SRC_BART = tt.data.Field(include_lengths=True, # include lengths
                             batch_first=True, # batches will be batch_size x
    ↳ max_len
                             tokenize=bart_tokenize, # use bart tokenizer
                             use_vocab=False, # bart tokenizer already
    ↳ converts to int ids
                             pad_token=bart_tokenizer.pad_token_id
                             )
TGT_BART = tt.data.Field(include_lengths=False,
                         batch_first=True, # batches will be batch_size x
    ↳ max_len
                         tokenize=bart_tokenize, # use bart tokenizer
                         use_vocab=False, # bart tokenizer already
    ↳ converts to int ids
                         pad_token=bart_tokenizer.pad_token_id
                         )
fields_bart = [('src', SRC_BART), ('tgt', TGT_BART)]

# Make splits for data
train_data_bart, val_data_bart, test_data_bart = tt.datasets.TranslationDataset.
    ↳ splits(
        ('_flightid.nl', '_flightid.sql'), fields_bart, path='./data/',
        train='train', validation='dev', test='test')

BATCH_SIZE = 1 # batch size for training/validation
```

```

TEST_BATCH_SIZE = 1 # batch size for test, we use 1 to make beam search
↳ implementation easier

train_iter_bart, val_iter_bart = tt.data.BucketIterator.
↳ splits((train_data_bart, val_data_bart),

                                                batch_size=BATCH_SIZE,
                                                device=device,
                                                repeat=False,
                                                sort_key=lambda x: len(x.
↳ src),

                                                sort_within_batch=True)

test_iter_bart = tt.data.BucketIterator(test_data_bart,
                                        batch_size=1,
                                        device=device,
                                        repeat=False,
                                        sort=False,
                                        train=False)

```

Let's take a look at the batch. Note that the shape of the batch is `batch_size x max_len`, instead of `max_len x batch_size` as in the previous part.

```

[ ]: batch = next(iter(train_iter_bart))
train_batch_text, train_batch_text_lengths = batch.src
print (f"Size of text batch: {train_batch_text.shape}")
print (f"First sentence in batch: {train_batch_text[0]}")
print (f"Length of the third sentence in batch: {train_batch_text_lengths[0]}")
print (f"Converted back to string: {bart_detokenize(train_batch_text[0])}")

train_batch_sql = batch.tgt
print (f"Size of sql batch: {train_batch_sql.shape}")
print (f"First sql in batch: {train_batch_sql[0]}")
print (f"Converted back to string: {bart_detokenize(train_batch_sql[0])}")

```

Now we are ready to implement the BART-based approach for the text-to-SQL conversion problem. In the below BART class, we have provided the constructor `__init__`, the `forward` function, and the `predict` function. Your job is to implement the main optimization `train_all`, and `evaluate_ppl` for evaluating validation perplexity for model selection.

Hint: you can use almost the same `train_all` and `evaluate_ppl` function you implemented before, but here a major difference is that due to setting `batch_first=True`, the batched source/target tensors are of size `batch_size x max_len`, as opposed to `max_len x batch_size` in the LSTM-based approach, and you need to make changes in `train_all` and `evaluate_ppl` accordingly.

```

[ ]: #TODO - finish implementing the `BART` class.
class BART(nn.Module):
    def __init__(self, tokenizer, pretrained_bart):
        """

```

```

Initializer. Creates network modules and loss function.
Arguments:
    tokenizer: BART tokenizer
    pretrained_bart: pretrained BART
"""
super(BART, self).__init__()

self.V_tgt = len(tokenizer)

# Get special word ids
self.padding_id_tgt = tokenizer.pad_token_id

# Create essential modules
self.bart = pretrained_bart

# Create loss function
self.loss_function = nn.CrossEntropyLoss(reduction="sum",
                                           ignore_index=self.padding_id_tgt)

def forward(self, src, src_lengths, tgt_in):
    """
    Performs forward computation, returns logits.
    Arguments:
        src: src batch of size (batch_size, max_src_len)
        src_lengths: src lengths of size (batch_size)
        tgt_in: a tensor of size (tgt_len, bsz)
    """
    # BART assumes inputs to be batch-first
    # This single function is forwarding both encoder and decoder (w/ cross_
    ↪ attn),
    # using `input_ids` as encoder inputs, and `decoder_input_ids`
    # as decoder inputs.
    logits = self.bart(input_ids=src,
                       decoder_input_ids=tgt_in,
                       use_cache=False
                       ).logits

    return logits

def evaluate_ppl(self, iterator):
    """Returns the model's perplexity on a given dataset `iterator`."""
    #TODO - implement this function
    ...
    ppl = ...
    return ppl

def train_all(self, train_iter, val_iter, epochs=10, learning_rate=0.001):
    """Train the model."""

```



```

#TODO - implement this function
...

def predict(self, tokens, K=1, max_T=400):
    """
    Generates the target sequence given the source sequence using beam search,
    ↳ decoding.
    Note that for simplicity, we only use batch size 1.
    Arguments:
        tokens: a list of strings, the source sentence.
        max_T: at most proceed this many steps of decoding
    Returns:
        a string of the generated target sentence.
    """
    string = ' '.join(tokens) # first convert to a string
    # Tokenize and map to a list of word ids
    inputs = torch.LongTensor(bart_tokenize(string)).to(device).view(1, -1)
    # The `transformers` package provides built-in beam search support
    prediction = self.bart.generate(inputs,
                                    num_beams=K,
                                    max_length=max_T,
                                    early_stopping=True,
                                    no_repeat_ngram_size=0,
                                    decoder_start_token_id=0,
                                    use_cache=True)[0]

    return bart_detokenize(prediction)

```

The code below will kick off training, and evaluate the validation perplexity. You should expect to see a value very close to 1.

```

[ ]: EPOCHS = 5 # epochs, we recommend starting with a smaller number like 1
      LEARNING_RATE = 1e-5 # learning rate

# Instantiate and train classifier
bart_model = BART(bart_tokenizer,
                  pretrained_bart
                  ).to(device)

bart_model.train_all(train_iter_bart, val_iter_bart, epochs=EPOCHS,
    ↳ learning_rate=LEARNING_RATE)
bart_model.load_state_dict(bart_model.best_model)

# Evaluate model performance, the expected value should be < 1.2
print (f'Validation perplexity: {bart_model.evaluate_ppl(val_iter_bart):.3f}')

```

As before, make sure that your model is making reasonable predictions on a few examples before evaluating on the entire test set.

```
[ ]: def bart_trial(sentence, gold_sql):
    print("Sentence: ", sentence, "\n")
    tokens = tokenize(sentence)

    predicted_sql = bart_model.predict(tokens, K=1, max_T=300)
    print("Predicted SQL:\n\n", predicted_sql, "\n")

    if verify(predicted_sql, gold_sql, silent=False):
        print ('Correct!')
    else:
        print ('Incorrect!')
```

```
[ ]: bart_trial(example_1, gold_sql_1)
```

```
[ ]: bart_trial(example_2, gold_sql_2)
```

```
[ ]: bart_trial(example_3, gold_sql_3)
```

```
[ ]: bart_trial(example_4, gold_sql_4)
```

```
[ ]: bart_trial(example_5, gold_sql_5)
```

```
[ ]: bart_trial(example_6, gold_sql_6b)
```

```
[ ]: bart_trial(example_7, gold_sql_7b)
```

```
[ ]: bart_trial(example_8, gold_sql_8)
```

### 6.2.1 Evaluation

The code below will evaluate on the entire test set. You should expect to see precision/recall/F1 greater than 40%.

```
[ ]: def seq2seq_predictor_bart(tokens):
    prediction = bart_model.predict(tokens, K=4, max_T=400)
    return prediction
```

```
[ ]: precision, recall, f1 = evaluate(seq2seq_predictor_bart, test_iter.dataset,
    ↪ num_examples=0)
print(f"precision: {precision:3.2f}")
print(f"recall:    {recall:3.2f}")
print(f"F1:        {f1:3.2f}")
```

## 7 Debrief

**Question:** We're interested in any thoughts you have about this project segment so that we can improve it for later years, and to inform later segments for this year. Please list any issues that

arose or comments you have to improve the project segment. Useful things to comment on might include the following:

- Was the project segment clear or unclear? Which portions?
- Were the readings appropriate background for the project segment?
- Are there additions or changes you think would make the project segment better?

but you should comment on whatever aspects you found especially positive or negative.

*Type your answer here, replacing this text.*

## 8 Instructions for submission of the project segment

This project segment should be submitted to Gradescope at <http://go.cs187.info/project4-submit-code> and <http://go.cs187.info/project4-submit-pdf>, which will be made available some time before the due date.

Project segment notebooks are manually graded, not autograded using otter as labs are. (Otter is used within project segment notebooks to synchronize distribution and solution code however.)

**We will not run your notebook before grading it.** Instead, we ask that you submit the already freshly run notebook. The best method is to “restart kernel and run all cells”, allowing time for all cells to be run to completion. You should submit your code to Gradescope at the code submission assignment at <http://go.cs187.info/project4-submit-code>. Make sure that you are also submitting your `data/grammar` file as part of your solution code as well.

We also request that you **submit a PDF of the freshly run notebook**. The simplest method is to use “Export notebook to PDF”, which will render the notebook to PDF via LaTeX. If that doesn’t work, the method that seems to be most reliable is to export the notebook as HTML (if you are using Jupyter Notebook, you can do so using `File -> Print Preview`), open the HTML in a browser, and print it to a file. Then make sure to add the file to your git commit. Please name the file the same name as this notebook, but with a `.pdf` extension. (Conveniently, the methods just described will use that name by default.) You can then perform a git commit and push and submit the commit to Gradescope at <http://go.cs187.info/project4-submit-pdf>.

## End of project segment 4